

Assignment

Comparing Methods on a Profit-Collecting VRPTW

Team Orienteering with Time Windows (TOPTW)

Applied AI: Academic perspectives: Combinatorial Optimisation Module

Goal. In this assignment you will compare *two* approaches among the techniques seen in class (Local Search, Simulated Annealing, Adaptive Large Neighbourhood Search, and, for small instances, exact methods). The emphasis is on the *process*: modelling choices, neighbourhood/operator design, parameter tuning, experimental method, and analysis of results. State-of-the-art optimality is *not* the goal.

1 Problem Summary (TOPTW)

You are given a depot and a set of customers with coordinates, demands, service times and *time windows*. Each customer i yields a profit p_i if served. The task is to select a subset of customers and partition them over at most K vehicle routes (or pay a fixed cost per used vehicle) such that *total collected profit minus routing/vehicle costs* is maximised, while respecting capacity and time-window feasibility per route.

This problem is commonly known as the *Team Orienteering Problem with Time Windows* (TOPTW), and can be seen as a profit-collecting variant of the VRPTW. It is a good fit for ruin-and-recreate methods like ALNS and remains compatible with VRPTW neighbourhoods used in Local Search and SA.

2 Mathematical Model (Hard and Soft Time Windows)

We give a compact MIP with time and load propagation. This VRPTW-style flow with time ordering eliminates subtours.

Data

- Depot 0; customers $V = \{1, \dots, n\}$; vehicles $k \in \{1, \dots, K\}$.
- Coordinates induce distances $d_{ij} \geq 0$ and travel times $t_{ij} \geq 0$ for $i, j \in \{0\} \cup V$, $i \neq j$ (assume $t_{ij} = d_{ij}$ unless stated).
- Demand $q_i \geq 0$, service time $s_i \geq 0$, time window $[a_i, b_i]$, profit $p_i \geq 0$ for $i \in V$.
- Vehicle capacity $Q > 0$; optional per-route duration limit $T_{\max} > 0$.
- Optional fixed cost $f \geq 0$ per used vehicle; lateness penalty $\beta \geq 0$ (only in the soft-TW variant).
- Big- M constants M_t, M_q sufficiently large to deactivate time/load propagation when an arc is unused.

Decision Variables

- $x_{ijk} \in \{0, 1\}$: arc $(i \rightarrow j)$ is used on route k ($i \neq j$, $i, j \in \{0\} \cup V$).
- $y_i \in \{0, 1\}$: customer i is served.
- $u_k \in \{0, 1\}$: vehicle k is used.
- $t_i \geq 0$: service start time at node i (one t_0 per route can be modelled via depot copies; below we use a single t_0 with big- M activation).
- $w_i \geq 0$: vehicle load after serving i .
- *Soft-TW only*: $L_i \geq 0$ lateness at i (penalised).

Objective

Hard time windows (no lateness):

$$\max \quad \sum_{i \in V} p_i y_i - \sum_{k=1}^K \sum_{\substack{i,j \in \{0\} \cup V \\ i \neq j}} d_{ij} x_{ijk} - f \sum_{k=1}^K u_k. \quad (1)$$

Soft time windows (lateness allowed, penalised):

$$\max \quad \sum_{i \in V} p_i y_i - \sum_{k=1}^K \sum_{\substack{i,j \in \{0\} \cup V \\ i \neq j}} d_{ij} x_{ijk} - f \sum_{k=1}^K u_k - \beta \sum_{i \in V} L_i. \quad (2)$$

Constraints

Vehicle usage and depot degree (per vehicle):

$$\sum_{j \in V} x_{0jk} = \sum_{i \in V} x_{i0k} = u_k \quad \forall k = 1, \dots, K. \quad (3)$$

Visit equals served (flow conservation over all vehicles):

$$\sum_{j \in \{0\} \cup V} \sum_{k=1}^K x_{ijk} = \sum_{j \in \{0\} \cup V} \sum_{k=1}^K x_{jik} = y_i \quad \forall i \in V. \quad (4)$$

Capacity propagation: For all $i \neq j$ in $\{0\} \cup V$,

$$w_j \geq w_i + q_j - M_q \left(1 - \sum_{k=1}^K x_{ijk} \right), \quad 0 \leq w_i \leq Q y_i \quad \forall i \in V. \quad (5)$$

Time propagation: For all $i \neq j$ in $\{0\} \cup V$,

$$t_j \geq t_i + s_i + t_{ij} - M_t \left(1 - \sum_{k=1}^K x_{ijk} \right). \quad (6)$$

Time windows: For all $i \in V$,

$$a_i \leq t_i \leq b_i + M_t(1 - y_i) \quad (\text{hard-TW}) \quad (7)$$

or, in the soft-TW variant,

$$a_i \leq t_i \leq b_i + L_i + M_t(1 - y_i), \quad L_i \geq 0. \quad (8)$$

Route duration (optional, if modelled with depot times):

$$t_0^{\text{return},k} \leq t_0^{\text{start},k} + T_{\max} \quad \forall k, \quad (9)$$

which can be implemented via depot copies per vehicle or accumulated travel/service-time counters.

Remark. The VRPTW-style time ordering together with degree/flow constraints removes subtours: time strictly progresses along used arcs, so cycles not connected to the depot are infeasible. Customer selection is governed by y_i . Calibrate M_t, M_q tightly (e.g. using bounds derived from b_i and Q) to improve LP relaxations.

3 Methods to Implement (Choose Two)

Pick any two of the following and ensure a fair experimental comparison. Make well considered choices!

Exact (LP/MIP or B&B). On small instances only (e.g. $n \leq 30\text{--}40$).

Local Search (LS).

Simulated Annealing (SA).

Adaptive Large Neighbourhood Search (ALNS).

4 Experimental Protocol

Instances. Use the curated TOP/TOPTW sets derived from Solomon VRPTW (see Appendix A). For a soft-TW variant, use the same coordinates/TWs but allow lateness with penalty β as in Appendix A.

Difficulty tiers.

- *Easy:* $n \approx 50$ (choose 3 instances; prefer class C or wider windows).
- *Hard:* $n \approx 100$ (choose 3 instances; include R/RC classes).

Stochastic fairness. For metaheuristics, run 5 independent seeds per instance. Report *mean*, *std*, and *best* objective.

Budgets. Same wall-clock budget per method and per instance (e.g. 2–5 minutes), identical initialisation policy across methods.

Metrics. Primary: assignment objective. Secondary: served customers, total distance, vehicles used, runtime.

Tuning. Use one development instance in each tier to tune a small parameter grid (e.g. SA cooling factor, ALNS scoring window, β). Fix parameters before final evaluation.

5 Deliverables

D1. Reproducible Jupyter notebook (Python) that:

- Parses instances (format in Appendix B).
- Implements your two chosen methods.
- Sets random seeds; logs runtime and objective.
- Produces summary tables/plots (per instance and aggregated).

D2. Short report (max 5 pages) covering:

- Modelling choices (hard vs soft TW, objective weights f, β).
- Neighbourhoods/operators and acceptance criteria.
- Tuning process (what you tried; how you selected final settings).
- Results with tables (*best/avg/std*), discussion of trade-offs.
- Recommendation: which method you would choose and why.

6 Practical Guidance

Initialisation. Greedy insertion by profit density $p_i/(d_{0i} + s_i)$ or regret-2 on the penalised objective.

Feasibility. Under hard TW, disallow infeasible moves. Under soft TW, allow temporary violations but penalise with βL_i (time-warp idea).

Relatedness (Shaw). Combine Euclidean distance, TW overlap, and profit similarity to select related nodes for removal.

Operator adaptation (ALNS). Typical scores: +5 for global improvement, +2 for best-in-iteration, +1 for accepted; update weights every fixed window (e.g. every 100 iterations).

A Benchmark Instances (Hard and Soft)

Hard time windows (public TOPTW)

Use the widely adopted TOP/TOPTW instances derived from Solomon (classes C/R/RC; sizes around 50 and 100). Choose:

- **Easy tier (3 instances):** $n \approx 50$; include at least one class C (clustered) and one non-C.
- **Hard tier (3 instances):** $n \approx 100$; include R and/or RC classes (typically harder).

Many repositories list for each instance the best-known or optimal values (for small sizes). In your report, cite the source you used and indicate whether you compare against optimal values or best-known solutions where available.

Soft-TW variant (derived)

To study penalty-based feasibility, derive a soft-TW version of any hard-TW instance by introducing nonnegative lateness variables L_i and adding $\beta \sum_i L_i$ to the objective. Replace the constraint $t_i \leq b_i$ by $t_i \leq b_i + L_i$. Recommended starting range for β : 5 to 10 (tune on your development instance). Keep all other data unchanged.

Note on difficulty. Class C (clustered) with wider windows is often easier; R/RC and tighter windows are harder. Larger n increases difficulty.

B Simple File Format (CSV + JSON)

Provide each instance with two files.

nodes.csv

```
id,type,x,y,profit,demand,service,tw_start,tw_end
0,depot,50.0,50.0,0,0,0,0,1e9
1,customer,12.3,77.1,84,3,10,120,240
...
...
```

meta.json

```
{
  "K": 3,
  "Q": 100,
  "vehicle_fixed_cost": 200.0,
  "speed": 1.0,
  "distance_cost": 1.0,
  "T_max": null,
  "time_windows": "hard",           // "hard" or "soft"
  "lateness_penalty_beta": 8.0     // used only if "soft"
}
```

Compute d_{ij} as Euclidean distance with double precision and set $t_{ij} = d_{ij}$ unless specified otherwise.

C Academic Integrity & Reproducibility

Your notebook must allow full reproduction: list exact instance names, random seeds, parameter settings, library versions, and measured wall-clock limits. Any external code you adapt must be clearly attributed.

D Suggested Reading (Optional)

Surveys and benchmark pages on Orienteering/TOPTW and Solomon-derived VRPTW data are widely available. Good starting points include a review of Orienteering variants, algorithmic studies on TOPTW (e.g. ILS/SA/ALNS), and the classic Solomon VRPTW description of instance classes (C/R/RC) and window patterns.